

Issued on: Sept 29, 2017

Due by 4 PM EST, Oct 07, 2017

Implement solution for this assignment on Cloudera Quick Start VM with CDH 5.12.

**Problem 1.** Download Quick Start VM for CDH 5.12 from [https://www.cloudera.com/downloads/quickstart\\_vms/5-8.html](https://www.cloudera.com/downloads/quickstart_vms/5-8.html). Start the VM. Please assign to the VM as much memory as you can. Examine whether `hadoop-hdfs-*`, `hadoop-mapreduce-*` and `hadoop-yarn-*` daemons are running. If those daemons are not running start all of them. If any of daemons fails to run, try to fix it.

[10%]

```
[[cloudera@quickstart conf.pseudo]$ for x in `cd /etc/init.d ; ls hadoop-*`; do sudo service $x status ; done
```

```
Hadoop datanode is running           [ OK ]
Hadoop journalnode is running        [ OK ]
Hadoop namenode is running           [ OK ]
Hadoop secondarynamenode is running  [ OK ]
Hadoop httpfs is running              [ OK ]
Hadoop historyserver is running       [ OK ]
Hadoop nodemanager is running         [ OK ]
Hadoop proxyserver is dead and pid file exists [ OK ]
Hadoop resourcemanager is running     [ OK ]
```

**Problem 2.** Examine whether there are HDFS directories for users: `spark`, `hive`, `oozie`, and `cloudera`. If the directories are present, find the content of those directories. If the directories are not present, create them. Please do not format the namenode.

[10%]

```
[cloudera@quickstart init.d]$ hadoop fs -ls /user
```

```
Found 8 items
```

```
drwxr-xr-x - cloudera cloudera      0 2017-07-19 06:28 /user/cloudera
drwxr-xr-x - mapred  hadoop          0 2017-07-19 06:29 /user/history
drwxrwxrwx - hive   supergroup      0 2017-07-19 06:31 /user/hive
drwxrwxrwx - hue    supergroup      0 2017-07-19 06:30 /user/hue
drwxrwxrwx - jenkins supergroup     0 2017-07-19 06:29 /user/jenkins
drwxrwxrwx - oozie  supergroup      0 2017-07-19 06:30 /user/oozie
drwxrwxrwx - root   supergroup      0 2017-07-19 06:29 /user/root
drwxr-xr-x - hdfs   supergroup      0 2017-07-19 06:31 /user/spark
```

```

[cloudera@quickstart init.d]$ hadoop fs -ls /user/spark
Found 1 items
drwxrwxrwx - spark supergroup      0 2017-10-02 12:18
/user/spark/applicationHistory
[cloudera@quickstart init.d]$ hadoop fs -ls /user/hive
Found 1 items
drwxrwxrwx - hive supergroup        0 2017-07-19 06:31 /user/hive/warehouse
[cloudera@quickstart init.d]$ hadoop fs -ls /user/oozie
Found 1 items
drwxrwxrwx - oozie supergroup        0 2017-07-19 06:30 /user/oozie/share
[cloudera@quickstart init.d]$ hadoop fs -ls /user/cloudera

```

**Problem 3.** Create new Linux user `smith`. Make that user a member of the `mapred` Linux group. Make that user a `sudo` user. Create the home directory of user `smith` in HDFS. Download provided files `bible.tar` and `shakespeare.tar`. Unzip both tar files and copy the resulting files into HDFS directory `input` of user `smith`. As user `smith` run Hadoop `grep` on both `bible` and `shakespeare` texts. Every Hadoop run requires separate output directory. Examine content of first 20 lines of files generated by Hadoop `grep`.

[15%]

user: smith

password: cloudera

## BIBLE OUTPUT

```

hadoop jar /usr/lib/hadoop-mapreduce/hadoop-mapreduce-examples.jar grep
/user/smith/bible/all-bible /user/smith/bible/all-bible-freq '\w+'

[smith@quickstart ~]$ hadoop fs -cat /user/smith/bible/all-bible-freq/part-r-00000 |
head -20
62394 the
38985 and
34654 of
13526 to
12846 And
12603 that
12445 in
9764 shall
9672 he
8940 unto
8854 I
8385 his
8057 a

```

```
7270 for
6974 they
6913 be
6884 is
6649 him
6647 LORD
6591 not
```

## Shakespeare Output

```
hadoop jar /usr/lib/hadoop-mapreduce/hadoop-mapreduce-examples.jar grep
/user/smith/shakespeare/all-shakespeare /user/smith/shakespeare/all-shakespeare-
freq '\w+'
```

```
[smith@quickstart ~]$ hadoop fs -cat /user/smith/shakespeare/all-shakespeare-
freq/part-r-00000 | head -20
```

```
25578 the
23027 I
19654 and
17462 to
16444 of
13524 a
12697 you
11296 my
10699 in
8857 is
8851 that
8402 not
8033 me
8020 s
7800 And
7231 with
7165 it
6812 his
6753 be
6246 your
```

**Problem 4.** Create your own version of “Hadoop grep” program using Spark. Compare your results with the results of Hadoop grep when applied to the texts of King James Bible, and all of Shakespeare’s works, contained in files `bible.tar` and `shakespear.tar` respectively. Notice small differences between results obtained by your Spark program and Hadoop grep. Try to explain what causes those differences. Save

results of your Spark grep operations both in HDFS and on your local file system. You can implement your solution using one of interactive shells or a standalone program. [20%]

## LOCAL FILE SYSTEM

### BIBLE

```
ul = sc.textFile("file:///home/cloudera/Desktop/all-bible")
counts = ul.flatMap(lambda x:x.split(" ")).map(lambda x:
(x,1)).reduceByKey(add)
```

```
print(counts.take(5))
```

```
exchanged = counts.map(lambda x: (x[1],x[0]))
print(exchanged.take(5))
sorted = exchanged.sortBy(lambda x: x[0], ascending=False)
print(sorted.take(20))
```

```
[(605968, u''), (62384, u'the'), (38711, u'and'), (34618, u'of'), (13505, u'to'), (12735,
u'And'), (12478, u'that'), (12279, u'in'), (9764, u'shall'), (9513, u'he'), (8930, u'unto'),
(8708, u'I'), (8362, u'his'), (8054, u'a'), (7183, u'for'), (6897, u'they'), (6754, u'be'), (6747,
u'is'), (6047, u'with'), (5878, u'not')]
```

### Shakespeare

```
shake_ul = sc.textFile("file:///home/cloudera/Desktop/all-shakespeare")
shake_counts = shake_ul.flatMap(lambda x:x.split(" ")).map(lambda x:
(x,1)).reduceByKey(add)
```

```
print(shake_counts.take(5))
```

```
shake_exchanged = shake_counts.map(lambda x: (x[1],x[0]))
print(shake_exchanged.take(5))
shake_sorted = shake_exchanged.sortBy(lambda x: x[0], ascending=False)
print(shake_sorted.take(20))
```

```
[(64531, u''), (25069, u'the'), (18793, u'and'), (16436, u'to'), (16069, u'of'), (15223, u'I'),
(12982, u'a'), (11180, u'my'), (10134, u'in'), (9109, u'you'), (8109, u'is'), (7773, u'that'),
(7123, u'not'), (7001, u'with'), (6594, u'his'), (6202, u'be'), (6119, u'your'), (5955,
u'tAnd'), (5781, u'for'), (5311, u'have')]
```

## Explanations

It looks like the hadoop grep removes the blank or empty spaces, which my spark grep did not. Also, I did not do .lower() in my spark grep script which also can account for my

differences in counts.

**Problem 5.** Create your own tables `KINGJAMES` with columns for words and frequencies and insert into the table the result of your Spark grep program which produces word counts in file `bible`. Find all words in table `KINGJAMES` which start with letter “w” and are 4 or more characters long and appear more than 250 times. Write a query that will tell us the number of such words. Before counting turn all words in lower case. When comparing a word with a string you use `LIKE` operator, like

```
word like 'a%' or word like '%th%'
```

Symbol `'%'` means any number of characters. You measure the length of a string using function `length()` and you change the case of a word to all lower characters using function `lower()`.

[20%]

```
bible_lines = sc.textFile("file:///home/cloudera/Desktop/all-bible")\
    .flatMap(lambda l: l.split())\
    .map(lambda x: re.sub("[^a-zA-Z]+", "", x.lower().encode("utf-8", "ignore")))\
    .filter(lambda x: x != "")

print(bible_lines.take(10))

bible_words = bible_lines.map(lambda p: Row( bible_word=str(p) ) )

print(bible_words.take(10))

bible_df = sqlContext.createDataFrame(bible_words)
print(bible_df.show(10))
bible_df.registerTempTable("KINGJAMES")

bible = sqlContext.sql("""
    SELECT
        bible_word,
        COUNT(*) freq
    FROM KINGJAMES
    WHERE lower(bible_word) like 'w%'
    AND length(bible_word) > 4
    GROUP BY bible_word
    HAVING COUNT(*) > 250
    """)

bible = bible.orderBy(bible['freq'].desc())
print(bible.show())
```

## OUTPUT

```
['king', 'james', 'bible', 'body', 'backgroundfaebd', 'margin',
'textalignjustify', 'p', 'textindent', 'em']
```

```
[Row(bible_word='king'), Row(bible_word='james'),
Row(bible_word='bible'), Row(bible_word='body'),
Row(bible_word='backgroundfaebd'), Row(bible_word='margin'),
Row(bible_word='textalignjustify'), Row(bible_word='p'),
Row(bible_word='textindent'), Row(bible_word='em')]
```

```
+-----+
|      bible_word|
+-----+
|           king|
|           james|
|           bible|
|           body|
| backgroundfaebd|
|           margin|
|textalignjustify|
|           p|
|       textindent|
|           em|
+-----+
```

only showing top 10 rows

None

```
+-----+-----+
|bible_word|freq|
+-----+-----+
|      which|4427|
|      words| 548|
|      would| 451|
| without| 442|
|      where| 407|
|      water| 396|
|      woman| 357|
| wherefore| 348|
|      wicked| 344|
|      whose| 314|
|wilderness| 304|
|      works| 302|
|      world| 287|
|      waters| 287|
|      written| 283|
+-----+-----+
```

**Problem 6.** Transfer content of your Hive KINGJAMES table to a Spark DataFrame. Perform the analysis from problem 6 using any available API in Spark. Please note that you are working with Spark 1.6.

[20%]

## CODE

```
hivecontext = HiveContext(sc)

dfs = hivecontext.sql("""
                        SELECT
                        freq,
                        lower(word) word
                        FROM kingjames
                        WHERE lower(word) like 'w%'
                        AND length(word) > 4
                        AND freq > 250
                        ORDER BY freq DESC
                        """)

print(dfs.show(20))
print(dfs.agg({"freq": "sum"}).show())
print(dfs.count())
```

## OUTPUT

```
+-----+-----+
|freq|      word|
+-----+-----+
|4297|    which|
| 546|    words|
| 443|    would|
| 436| without|
| 396|    water|
| 355|    woman|
| 343|   wicked|
| 335|    where|
| 304|wilderness|
| 301|    works|
| 288|    world|
| 286|   waters|
| 284|    whose|
| 283|   written|
| 261| wherefore|
+-----+-----+

# CUMSUM of words in query
+-----+
|sum(freq)|
+-----+
|      9158|
+-----+
```

<b># Number of words in Query</b> <b>15</b>
--

**Problem 7.** Use Sqoop to transfer the content of MySQL database `retail_db` which is present on the Cloudera VM into Hive. Demonstrate that new Hive tables are created and correspond to the original MySQL tables. Find the number of rows in each table. Compare those row counts with row counts in MySQL database.  
[15%]

## MYSQL TABLES

```
mysql> show databases;
+-----+
| Database          |
+-----+
| information_schema |
| retail_db         |
+-----+
2 rows in set (0.01 sec)

mysql> use retail_db;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Database changed
mysql> show tables;
+-----+
| Tables_in_retail_db |
+-----+
| categories          |
| customers            |
| departments          |
| order_items          |
| orders               |
| products             |
+-----+
6 rows in set (0.00 sec)
```

## HIVE IMPORT

```
sqoop import-all-tables -m 1 --connect
jdbc:mysql://quickstart:3306/retail_db --username=retail_dba --
password=cloudera --compression-codec=snappy --as-parquetfile --
warehouse-dir=/user/hive/warehouse --hive-import
```



```
$> beeline
!connect jdbc:hive2://127.0.0.1:10000/default hive cloudera org.apache.hive.jdbc.HiveDriver
```

## HIVE TABLES

```
0: jdbc:hive2://127.0.0.1:10000/default> show tables;
INFO  : Compiling command(queryId=hive_20171006161717_465c7d95-37a2-4cdc-9805-c32b28034e82): show tables
INFO  : Semantic Analysis Completed
INFO  : Returning Hive schema:
Schema(fieldSchemas:[FieldSchema(name:tab_name, type:string, comment:from deserializer)], properties:null)
INFO  : Completed compiling
command(queryId=hive_20171006161717_465c7d95-37a2-4cdc-9805-c32b28034e82); Time taken: 0.126 seconds
INFO  : Concurrency mode is disabled, not creating a lock manager
INFO  : Executing command(queryId=hive_20171006161717_465c7d95-37a2-4cdc-9805-c32b28034e82): show tables
INFO  : Starting task [Stage-0:DDL] in serial mode
INFO  : Completed executing
command(queryId=hive_20171006161717_465c7d95-37a2-4cdc-9805-c32b28034e82); Time taken: 0.018 seconds
INFO  : OK
+-----+---+
|  tab_name  |
+-----+---+
| categories |
| customers  |
| departments|
| kingjames  |
| order_items|
| orders     |
| products   |
+-----+---+
7 rows selected (0.256 seconds)
```

**Find the number of rows in each table. Compare those row counts with row counts in MySQL database.**

```
Categories
mysql> select count(*) from categories;
+-----+
| count(*) |
+-----+
|         58 |
+-----+
1 row in set (0.00 sec)
```

```
mysql> select count(*) from customers;
```

```
+-----+
```

```
| count(*) |
```

```
+-----+
```

```
|    12435 |
```

```
+-----+
```

```
1 row in set (0.04 sec)
```

```
0: jdbc:hive2://127.0.0.1:10000/default> select count(*) from  
categories;
```

```
+-----+
```

```
| _c0 |
```

```
+-----+
```

```
| 58 |
```

```
+-----+
```

```
1 row selected (25.214 seconds)
```

### Departments

```
mysql> select count(*) from departments;
```

```
+-----+
```

```
| count(*) |
```

```
+-----+
```

```
|         6 |
```

```
+-----+
```

```
1 row in set (0.00 sec)
```

```
0: jdbc:hive2://127.0.0.1:10000/default> select count(*) from  
departments;
```

```
+-----+
```

```
| _c0 |
```

```
+-----+
```

```
| 6 |
```

```
+-----+
```

```
1 row selected (24.858 seconds)
```

### Order Items

```
mysql> select count(*) from order_items;
```

```
+-----+
```

```
| count(*) |
```

```
+-----+
```

```
|    172198 |
```

```
+-----+
```

```
1 row in set (0.07 sec)
```

```
0: jdbc:hive2://127.0.0.1:10000/default> select count(*) from  
order_items;
```

```

+-----+---+
|  _c0  |
+-----+---+
| 172198 |
+-----+---+
1 row selected (25.533 seconds)

```

### Orders

```
mysql> select count(*) from orders;
```

```

+-----+---+
| count(*) |
+-----+---+
|      68883 |
+-----+---+
1 row in set (0.03 sec)

```

```
0: jdbc:hive2://127.0.0.1:10000/default> select count(*) from
orders;
```

```

+-----+---+
|  _c0  |
+-----+---+
| 68883 |
+-----+---+
1 row selected (26.502 seconds)

```

### Products

```
mysql> select count(*) from products;
```

```

+-----+---+
| count(*) |
+-----+---+
|      1345 |
+-----+---+
1 row in set (0.00 sec)

```

```
0: jdbc:hive2://127.0.0.1:10000/default> select count(*) from
products;
```

```

+-----+---+
|  _c0  |
+-----+---+
| 1345 |
+-----+---+
1 row selected (24.733 seconds)

```

We are also attaching two groups of example data files for Hive:  
`examples_older.zip` and `hive_examples.zip`. You might find those files useful if you want to keep on learning about the technology. You could get those files by downloading Hive distributions, as described in notes.

.