

2. Create a table for 5 languages in which you will collect statistics about the languages used.

```
: languages = [
    'English-Latin1',
    'Danish_Dansk-Latin1',
    'German_Deutsch-Latin1',
    'Filipino_Tagalog-Latin1',
    'Italian-Latin1'
]
```

3. Place in that table the number of words in each language in UDHR, number of unique words, average length of words, number of sentences contained in UDHR and average number of words per sentence.

```
##https://stackoverflow.com/questions/35900029/average-sentence-length-for-every-text-in-corpus-python3-nltk

np_object=[]
for lang in languages:
    chars_count = len(udhr.raw(lang))
    word_count = len(udhr.words(lang))
    unique_word_count = len(set(udhr.words(lang)))
    word_length_avg = round(chars_count/word_count)
    #word_length_avg = sum(len(sent) for sent in udhr.sents(fileids=[lang])) / len(udhr.sents(fileids=[lang]))
    sents_count = len(udhr.sents(lang))
    avg_num_words_per_sents = round(word_count/sents_count)
    x = [
        lang,
        word_count,
        unique_word_count,
        word_length_avg,
        sents_count,
        avg_num_words_per_sents
    ]
    np_object.append(x)
    x = []
df = pd.DataFrame(np_object, columns=['language', 'word_count', 'word_count_unique', 'word_length_avg', 'sents_count',
df
```

	language	word_count	word_count_unique	word_length_avg	sents_count	avg_num_words_per_sents
0	English-Latin1	1781	533	5.0	67	26.0
1	Danish_Dansk-Latin1	1696	584	5.0	86	19.0
2	German_Deutsch-Latin1	1521	579	6.0	60	25.0
3	Filipino_Tagalog-Latin1	1803	480	5.0	75	24.0
4	Italian-Latin1	1723	578	5.0	51	33.0

4. Create a distribution of sentence lengths for each language.

```
cfd = nltk.ConditionalFreqDist(
    (lang, len(sent))
    for lang in languages
    for sent in udhr.sents(lang))

# (lang, (sum(len(sent))/ len(udhr.sents(fileids=[lang])))
cfd


ConditionalFreqDist(nltk.probability.FreqDist,
    {'Danish_Dansk-Latin1': FreqDist({3: 28,
        6: 2,
        7: 3,
        10: 1,
        11: 3,
        12: 3,
        13: 1,
        14: 4,
        15: 3,
        16: 3,
        17: 1,
        18: 3,
        19: 3,
        20: 1,
        21: 2,
        22: 1,
        23: 2,
        26: 4,
        27: 1,
```

```
: cfd.tabulate(conditions=languages, samples=range(10), cumulative=True)
```

	0	1	2	3	4	5	6	7	8	9
English-Latin1	0	0	0	0	2	2	3	3	3	4
Danish_Dansk-Latin1	0	0	0	28	28	28	30	33	33	33
German_Deutsch-Latin1	0	0	0	1	1	2	2	2	3	5
Filipino_Tagalog-Latin1	0	0	13	13	23	23	23	23	23	24
Italian-Latin1	0	0	0	0	0	0	0	0	1	1


5. Plot those (non-cumulative) distributions on one diagram.

```
cfid.plot(cumulative=False)
```



Plotting for fun cumuallive distributions on one diagram

```
cfid.plot(cumulative=True)
```



Problem 2. Identify 10 most frequently used words longer than 7 characters in the entire corpus of Inaugural Addresses. Do not identify 10 words for every speech but rather 10 words for the entire corpus. Which among those words has the largest number of synonyms? List all

synonyms for those 10 words. Which one of those 10 words has the largest number of hyponyms? List all hyponyms of those 10 most frequently used “long” words. The purpose of this problem is to familiarize you with WordNet and concepts of synonyms and hyponyms.

(25%)

Your literature for Problems 1 and 2 are chapters 1 and 2 of Natural Language Processing with Python book by Steven Bird et al.

See p2.ipyb for a closer look

1. Identify 10 most frequently used words longer than 7 characters in the entire corpus of Inaugural Addresses:

- Do not identify 10 words for every speech but rather 10 words for the entire corpus.

```
word_extract=[word.lower() for word in inaugural_corpus if (len(word) > 7 and word.isalpha())]
word_frequency=FreqDist(word_extract)
print("\n Top 10 words in Inauguaral Corpus-- ignores case and non-alpha chars")
word_frequency.tabulate(10)
print word_freq
```

```
Top 10 words in Inauguaral Corpus-- ignores case and non-alpha chars
government      citizens constitution      national      american      congress      interests      political      execu
principles
593             237             205             154             147             129             113             106
93
None
```

2. Which among those words has the largest number of synonyms?

- List all synonyms for those 10 words.

```
# wn.synsets('motorcar')
top_10_words = [
    'government',
    'citizens',
    'constitution',
    'national',
    'american',
    'congress',
    'interests',
    'political',
    'executive',
    'principles'
]
syn_words = {}
for word in top_10_words:
    syn_words[word] = []
    syn = wn.synsets(word)
    for s in syn:
        lemmas = s.lemma_names()
        if len(lemmas) > 0:
            for l in lemmas:
                l = l.lower()
                if l is not word and l.encode('utf-8') not in syn_words[word]:
                    syn_words[word].append(l.encode('utf-8'))
pprint.pprint(syn_words)

# pprint.pprint([i for synset in wn.synsets(word) for i in synset.lemma_names()])
```

```
{
  'american': ['american', 'american_english', 'american_language'],
  'citizens': ['citizen'],
  'congress': [
    'congress',
    'united_states_congress',
    'u.s._congress',
    'us_congress',
    'sexual_intercourse',
    'intercourse',
    'sex_act',
    'copulation',
    'coitus',
    'coition',
    'sexual_congress',
    'sexual_relation',
    'relation',
    'carnal_knowledge'
  ],
  'constitution': [
    'fundamental_law',
    'organic_law',
    'constitution',
    'establishment',
    'formation',
    'organization',
    'organisation',
    'united_states_constitution',
    'u.s._constitution',
    'us_constitution',
    'constitution_of_the_united_states',
    'composition',
    'physical_composition',
    'makeup',
    'make-up',
    'old_ironsides'
  ],
}
```

We can see that constitution is the top 10 word with the most synonyms

```
for word in syn_words:
    print(str(word)+" : " +str(len(syn_words[word])))
```

```
interests:12
executive:3
constitution:17
congress:15
government:9
national:5
citizens:1
political:1
principles:4
american:3
```

3. List all hyponyms of those 10 most frequently used “long” words. ¶

- Which one of those 10 words has the largest number of hyponyms?
- The purpose of this problem is to familiarize you with WordNet and concepts of synonyms and hyponyms.

```
] top_10_words = [
    'government',
    'citizens',
    'constitution',
    'national',
    'american',
    'congress',
    'interests',
    'political',
    'executive',
    'principles'
]
hyponym_words = {}
for word in top_10_words:
    hyponym_words[word] = []
    syn = wn.synsets(word)
    for s in syn:
        hypos = s.hyponyms()
        if len(hypos) > 0:
            for hyp in hypos:
                for h in hyp.lemma_names():
                    h = h.lower()
                    if h is not word and h.encode('utf-8') not in hyponym_words[word]:
                        hyponym_words[word].append(h.encode('utf-8'))
pprint.pprint(hyponym_words)
```


FF--FF--FF--FF--FF--FF--FF--FF--

```
{ 'american': [ 'african-american',
                'african_american',
                'afro-american',
                'black_american',
                'alabaman',
                'alabamian',
                'alaskan',
                'anglo-american',
                'appalachian',
                'arizonan',
                'arizonian',
                'arkansan',
                'arkansawyer',
                'asian_american',
                'bay_stater',
                'bostonian',
                'californian',
                'carolinian',
                'coloradan',
                'connecticutan'
```

We can see that american is the top 10 word with the most hyponyms

```
for word in hyponym_words:
    print(str(word)+": " +str(len(hyponym_words[word])))
```

```
interests:43
executive:18
constitution:17
congress:18
government:32
national:4
citizens:9
political:0
principles:62
american:109
```

**Problem 3. Create your own grammar for the following sentence:
“Describe every step of your work and present all intermediate and
final results in a Word document”.**

(10%)

**Your literature for Problem 3 is chapter 8 of Natural Language
Processing with Python book by Steven Bird et al.**

See p3.ipyb for a closer look

```
In [1]: import nltk
```

Problem 3

1. Create your own grammar for the following sentence:

- "Describe every step of your work and present all intermediate and final results in a Word document".

```
In [2]: sentence = """Describe every step of your work and present all intermediate and final results in a Word document"""
```

```
In [3]: tokens = nltk.word_tokenize(sentence)
```

```
In [4]: tokens
```

```
Out[4]: ['Describe',  
         'every',  
         'step',  
         'of',  
         'your',  
         'work',  
         'and',  
         'present',  
         'all',  
         'intermediate',  
         'and',  
         'final',  
         'results',  
         'in',  
         'a',  
         'Word',  
         'document']
```

```
In [5]: tagged_text = nltk.pos_tag(tokens)
```

```
In [6]: tagged_text
```

```
Out[6]: [('Describe', 'NNP'),  
         ('every', 'DT'),  
         ('step', 'NN'),  
         ('of', 'IN'),  
         ('your', 'PRP$'),  
         ('work', 'NN'),  
         ('and', 'CC'),  
         ('present', 'JJ'),  
         ('all', 'DT'),  
         ('intermediate', 'JJ'),  
         ('and', 'CC'),  
         ('final', 'JJ'),  
         ('results', 'NNS'),  
         ('in', 'IN'),  
         ('a', 'DT'),  
         ('Word', 'NNP'),  
         ('document', 'NN')]
```

Problem 4. Install and compile Word2Vec C executables. Train CBOW model and create 200 dimensional embedding of Word Vectors. Demonstrate that you could run analogical reasoning when searching for country's favorite food starting with japan and sushi. Note that words might have to be in lower case. Find favorite food for 5 different countries. Report improbable results as well as good results. Use scripts provided with original Google C code.
(20%)

```

Enter word or sentence (EXIT to break): EXIT
swaite@Rmt-mac-swaite:~/stirling/CSIE-63/assignment-8/word2vec|master$
⇒ sudo ./demo-analogy.sh
make: Nothing to be done for `all'.

```

Note that for the word analogy to perform well, the models should be trained on much larger data set
 Example input: paris france berlin

```

Starting training using file text8
Vocab size: 71290
Words in train file: 16718843
Alpha: 0.000123 Progress: 99.57% Words/thread/sec: 73.12k
real    0m43.345s
user    3m59.424s
sys     0m1.064s

```

japan sushi thailand

Enter three words (EXIT to break): japan sushi thailand		
Word: japan	Position in vocabulary: 582	
Word: sushi	Position in vocabulary: 30679	
Word: thailand	Position in vocabulary: 5640	
Distance	Word	
	crab	0.561080
	bento	0.559273
	kimchi	0.541000
	mochi	0.540748
	crepe	0.536611

japan sushi italy

Enter three words (EXIT to break): japan sushi italy		
Word: japan	Position in vocabulary: 582	
Word: sushi	Position in vocabulary: 30679	
Word: italy	Position in vocabulary: 843	
Distance	Word	
	strawberries	0.500201
	kelp	0.500120
	sprouts	0.495927
	omelette	0.493959
	mussel	0.491622
	pies	0.489951
	cranberries	0.474655

japan sushi mexico

Enter three words (EXIT to break): japan sushi mexico

Word: japan Position in vocabulary: 582

Word: sushi Position in vocabulary: 30679

Word: mexico Position in vocabulary: 1352

Word	Distance
barbecued	0.556577
manioc	0.536024
tofu	0.514035
perforated	0.512456
kare	0.512113
soups	0.510998
leftover	0.506907
thud	0.506563
feeds	0.506198
cookbookwiki	0.503084
crispy	0.498741
boutiques	0.497999
subgenera	0.497614
leaved	0.496927
marinated	0.496811
grasshoppers	0.495767
snack	0.495525
glutinous	0.495342
seasonings	0.495112
chrysalis	0.494900

japan sushi canada

Enter three words (EXIT to break): japan sushi canada

Word: japan Position in vocabulary: 582

Word: sushi Position in vocabulary: 30679

Word: canada Position in vocabulary: 474

Word	Distance
kare	0.616777
glutinous	0.540952
cookbookwiki	0.535959
thyme	0.534985
boutiques	0.530151
hives	0.530043
lamiales	0.523256
caterpillar	0.521952
juniper	0.519933
cherry	0.518585
allo	0.518086
buckwheat	0.516296
sprouts	0.514408
surfing	0.513681
bullwinkle	0.512200
leftover	0.510922
bento	0.510452
sapphire	0.509470
pickles	0.509440
mussel	0.508209
molluscs	0.506993
roundworm	0.506693
paralel	0.505070

japan sushi india

0.493700
 Enter three words (EXIT to break): japan sushi india
 Word: japan Position in vocabulary: 582
 Word: sushi Position in vocabulary: 30679
 Word: india Position in vocabulary: 508

Word	Distance
-----	-----
cookbookwiki	0.534107
manioc	0.525004
juniper	0.514114
weed	0.513496
crispy	0.511961
mochi	0.511274
glutinous	0.509846
breaded	0.508894
drawl	0.506393
leftover	0.506105
sprouts	0.504681
shakin	0.504605
thumbnail	0.503094
cayenne	0.499609
lha	0.497450
soya	0.495425
parsley	0.495213
thud	0.494816
lamiales	0.494529
kare	0.494128
ebay	0.493733
food	0.493700

Problem 5. Install and run Genism Python Word2Vec API. Find the most probable words you will obtain when you start with an emperor add a woman and subtract a man. Use this tutorial as a guide

<https://rare-technologies.com/word2vec-tutorial/>
(20%)

See p5.ipynb for a closer look

```
In [1]: # import modules & set up logging
import gensim, logging
```

Working on a Macbook Pro so needed to use the link below to get this code work properly

- <https://github.com/William-Yeh/word2vec-mac>

```
In [6]: logging.basicConfig(format='%(asctime)s : %(levelname)s : %(message)s', level=logging.INFO)
```

```
In [7]: #load model, model was created in problem 1
model = gensim.models.KeyedVectors.load_word2vec_format('./word2vec/vectors.bin', binary=True)

2017-10-28 01:43:49,221 : INFO : loading projection weights from ./word2vec/vectors.bin
2017-10-28 01:43:50,188 : INFO : loaded (71290, 200) matrix from ./word2vec/vectors.bin
```

```
In [8]: #what requested in the problem
model.most_similar(positive=['emperor', 'woman'], negative=['man'], topn=1)

2017-10-28 01:43:50,236 : INFO : precomputing L2-norms of word weight vectors

Out[8]: [(u'montoku', 0.6238285303115845)]
```

Analysis of output

These results look correct. Looking up "montoku" on wikipedia gave the following results. https://en.wikipedia.org/wiki/Emperor_Montoku

"Emperor Montoku was the 55th emperor of Japan, according to the traditional order of succession. Montoku's reign lasted from 850 to 858."

Please, describe every step of your work and present all intermediate and final results in a Word document. Please, copy past text version of all essential command and snippets of results into the Word document with explanations of the purpose of those commands. We cannot retype text that is in JPG images. Please, always submit a separate copy of the original, working scripts and/or class files you used. Sometimes we need to run your code and retyping is too costly. Please include in your MS Word document only relevant portions of the console output or output files. Sometime either console output or the result file is too long and including it into the MS Word document makes that document too hard to read. PLEASE DO NOT EMBED files into your MS Word document. For issues and comments visit the class Discussion Board. If you use some other language other than Python in your daily work with NLP, please be free to use that language and a framework of your choice to do this assignment.